Hermes Agent

A Practical Guide for Founders & Operators

What it is, how it works, and whether it's worth your time

27,000+	14	$5/mo	6,000+
GitHub stars in 2 months	Messaging platforms	Minimum server cost	MCP integrations

Based on Hermes Agent v0.7.0 · April 2026 · MIT Open Source · Nous Research

Contents

01 What Hermes Actually Is

Most AI tools have a short memory and a shorter attention span. You open a chat, give it context, get an answer, close the window. Next time, you start from scratch. Every session is a first date.

Hermes Agent is built around a different premise: an AI that runs continuously in the background, remembers everything relevant across sessions, gets better at your specific workflows the more you use it, and can be reached from your phone, Slack, or Discord — not just a chat window on your laptop.

Nous Research, a small open-source AI lab, released it in February 2026. It reached 27,000 GitHub stars in two months — almost entirely from developers and operators who recognized it solved a real structural problem, not just a convenience one.

The Problem It's Solving

If you've used Claude, ChatGPT, or any AI assistant seriously, you've probably hit one or more of these walls:

- **The context wall:** You have to re-explain your business, preferences, and project every single conversation.
- **The availability wall:** The AI is only as useful as your laptop being open and you being present.
- **The consistency wall:** It gives great advice once, but forgets it by next week — so the same mistakes recur.
- **The maintenance wall:** You want it to follow your specific standards, but encoding those requires constant upkeep.

The Core Idea: A Self-Improving Harness

The key concept behind Hermes is something its builders call a **self-improving harness** — a set of instructions, constraints, and learned behaviors that the AI builds and maintains *for itself*, rather than requiring you to configure everything upfront.

Think of hiring a new contractor. At first, you explain everything: your preferences, your standards, what you've tried before. A mediocre contractor needs that re-explained every engagement. A great one builds their own notes, refines their approach based on your feedback, and shows up to the next project already knowing what matters to you. Hermes is built to be the second kind — except it never gets pulled to another client, works around the clock, and costs around $5–10 a month to run.

Key Numbers at a Glance

Metric	Figure
GitHub stars (2 months post-launch)	27,000+
Built-in tools	40+
External app integrations (via MCP)	6,000+
Platforms you can reach it from	14
Minimum monthly server cost	~$5
Memory footprint (no local AI model)	Under 500MB
License	MIT — fully open source

02 How It Learns: The Engine Under the Hood

You don't need to understand the technical internals. But understanding the logic of how Hermes learns will help you use it better and set realistic expectations. Three things work together: memory, Skills, and a learning loop that connects them.

Memory: Three Layers

Layer	What It Stores	Why It Matters
Session Memory (What just happened)	Every conversation, indexed and stored locally. Retrieved on demand, not bulk-loaded.	Even after months of use, it loads only what's relevant. Cost and speed stay constant.
Persistent Memory (Who you are)	Durable facts: your preferences, decisions, processes, what you've tried and ruled out.	Come back after two weeks, say "continue that vendor research" — it knows exactly what you mean.
Skill Memory (How to do things)	Documented procedures for recurring tasks: triggers, steps, your specific preferences.	Recurring tasks run correctly without re-instruction. Approaches improve with every use.

PRIVACY NOTE — All memory data lives in a local file on your server — not in Nous Research's cloud. You own it, can back it up, and take it with you if you ever switch servers.

Skills: Your SOPs, Built Automatically

In Hermes, a Skill is a plain-text document capturing how to handle a specific recurring task. It contains the trigger (when to use it), the steps to follow, your preferences, and things to watch out for.

Concrete example: You ask Hermes every Monday to pull weekend sales data and send a summary. After a few iterations where you say "also flag anything where refunds exceeded 10% of orders," Hermes doesn't just do that once — it edits the Skill so that flagging refund anomalies is a permanent step. By week six, the Monday summary is exactly what you need without any guidance.

The Learning Loop

These three components connect in a continuous loop. You complete a task, Hermes writes relevant facts to memory, distills the approach into a Skill if it's complex and recurring, and incorporates your feedback automatically. The loop runs without you triggering it.

WHAT THIS MEANS IN PRACTICE — The first week feels like any other AI assistant — useful but not remarkable. By week four, it handles recurring tasks with noticeably less guidance. By month three, the gap between "Hermes knows how we work" and "this new tool has no idea about us" becomes significant.

03 What It Can Do: Tools, Integrations & Platforms

Built-In Tools: 40+ Capabilities Out of the Box

You won't use all of them — but knowing what exists helps you recognize when Hermes can handle something you'd otherwise do manually.

Category	What You Can Do
Execution (terminal, code sandbox, files)	Run commands, execute scripts in isolation, read and write files
Information (web, browser, history search)	Research topics, scrape pages, recall context from past conversations
Media (vision, image gen, text-to-speech)	Analyze screenshots, create visuals, produce audio summaries
Memory & Tasks (scheduler, Skill manager)	Read/write memory, manage Skills, set up recurring scheduled jobs
Coordination (sub-agents, multi-model)	Run parallel tasks, cross-check answers with multiple AI models simultaneously

Three Tools Worth Highlighting for Operators

- **Scheduled jobs:** Tell Hermes in plain English to do something on a schedule — "check our GitHub for new issues every morning and send me a summary" — and it sets up the schedule itself. No technical configuration required.
- **Sub-agent delegation:** Hermes can spawn up to three parallel agents to work simultaneously on different parts of a task. Research three vendors at once, get consolidated results in roughly half the time.
- **Session search:** When you reference something from a past conversation ("that supplier we discussed last month"), it searches stored history and retrieves the relevant context automatically.

Connecting to Your Existing Tools (MCP)

MCP (Model Context Protocol) lets Hermes plug into external services without any custom code. The ecosystem covers 6,000+ applications. Once connected, you interact with these tools through Hermes in plain language.

Integration	What Hermes Can Do
GitHub	Review PRs, create issues, track commits, generate commit reports
Databases (Postgres, MySQL, SQLite)	Query data in plain language, generate reports, spot trends — no SQL needed
Slack / Discord	Receive questions, send updates, participate in channels autonomously
Jira / Notion / Google Drive	Read and update project management tools and documents
Email / SMS	Send automated summaries, alerts, and scheduled reports

IMPORTANT — Database integrations have read and write access by default. If connecting to live business data, use a read-only database account. The capability to accidentally modify data is real.

Reach It From Anywhere: 14 Platforms

The critical design point: all platforms share the same brain. A conversation started on Telegram can be continued from Slack. Context never resets when you switch channels. There is one Hermes instance, one memory, one set of Skills — just multiple doors into it.

Platform	Best For	Setup Time
Telegram Bot	Personal access, mobile — most common starting point	Under 5 minutes
Discord	Team use, channel-based workflows	Under 10 minutes
Slack	Enterprise/team environments	Longer (App setup required)
CLI (terminal)	Direct server access	None — included by default
Email / SMS	Automated alerts and scheduled outputs	Under 10 minutes

04 Getting It Running

Three Deployment Options

Option	Best For	Cost	Notes
Local (your computer)	Testing before committing	$0	Only runs when your computer is on
$5 VPS (recommended)	Actual daily use, 24/7 uptime	$4–5/mo	Hetzner, DigitalOcean, or Vultr on Ubuntu 22.04
Serverless	Intermittent use, minimal idle cost	Near $0 idle	Wakes on incoming messages via Daytona or Modal

One-Line Install

On macOS, Linux, or a VPS, open a terminal and run:

```
curl -fsSL
https://raw.githubusercontent.com/NousResearch/hermes-agent/main/scripts/install.sh | bash
```

Type `hermes` to launch. If you see the welcome message, you're good to go.

Configuration: One File Handles Everything

All setup lives in `~/.hermes/config.yaml`. A working setup requires just a few fields:

Setting	What It Does	Recommended Value
provider	Which AI provider to use	openrouter — gives access to 200+ models, flexible switching
model	Which AI model to run	anthropic/claude-sonnet-4 or hermes-3-70b via Nous Portal
api_key	Your authentication key	From OpenRouter, Anthropic, or your chosen provider
telegram.token	Enables Telegram Bot access	From @BotFather in Telegram (under 5 min to get)

APRIL 2026 NOTE — Anthropic has restricted Claude access for third-party tools through Pro/Max subscriptions. You can still use Claude via direct API access (pay-per-use), but costs are higher. OpenRouter with Claude Sonnet is a practical default.

Your Data Stays With You

All of Hermes's state — conversation history, memory, Skills, config — lives in one directory on whatever machine you deploy it on. To back it up: copy the folder. To migrate: move the folder. Nothing is stored in anyone else's cloud.

05 Real-World Use Cases for Operators

Knowledge Management & Research

The typical pattern without Hermes: you research something across multiple sessions, spreading insights across different chat windows and notes. Return two weeks later and you're re-explaining everything from scratch.

With Hermes, multi-week research projects have continuity. You establish what you're investigating, what you've ruled out, and what you're prioritizing. Hermes retains all of that. Sub-agent delegation makes deep research faster — instead of researching three vendors sequentially over an afternoon, you dispatch three parallel agents and get consolidated findings in 20 minutes.

Operational Monitoring & Reporting

Hermes can monitor things on your behalf and push information to you proactively, rather than waiting for you to ask:

- **Daily briefings:** Receive a summary of GitHub activity, new support tickets, or overnight revenue data every morning at 8am.
- **Weekly reports:** Auto-generated based on real data pulled from connected tools, including context from the week's conversations and decisions.
- **Exception alerts:** "If refund rate in the past 24 hours exceeds X%, send me a Telegram message immediately."
- **Pipeline tracking:** Regular check-ins on deal status, automatically pulling from your CRM and summarizing movement.

Content & Communications Workflows

For founders who produce regular content — newsletters, investor updates, reports — Hermes's accumulating style memory is genuinely useful. Every correction you make (changing formal phrasing to casual, removing a cliché) gets incorporated into the style Skill. After six weeks, it has a detailed model of your voice that it applies automatically.

For content series (investor updates, weekly newsletters), Hermes maintains continuity between installments: what was covered last time, what threads are open, what the reader already knows.

Team Coordination & Async Communication

Deploying Hermes on Slack or Discord gives your team a shared AI resource with persistent memory of your business context. A team member can ask it to review a PR, pull a data summary, or draft a response to a complex customer situation — and it already knows your product, policies, and the relevant context from past decisions.

KEY POINT — Because memory is shared across platforms and users, Hermes builds a compounding understanding of how your whole team works — not just how one person uses it.

06 Hermes vs. Your Current AI Tools

There are now three meaningfully different categories of AI tools. They aren't competitors — they're different tools for different jobs.

Category	Examples	Core Strength
Interactive Assistants	Claude, ChatGPT	Tasks requiring your judgment throughout — writing, analysis, real-time decisions
Configuration-Driven Agents	OpenClaw	Enterprise scenarios where auditability, transparency, and consistency matter most
Autonomous Background Agents	Hermes	Monitoring, recurring workflows, and anything benefiting from continuity across weeks

When to Use Which Tool

Scenario	Best Tool	Why
Writing, analysis, product decisions	Claude / ChatGPT	Real-time feedback loop needed
Standardized team processes, compliance	OpenClaw	Transparent, auditable, predictable
24/7 monitoring, scheduled reports	Hermes	Runs unattended with cron scheduling
Multi-week research projects	Hermes	Cross-session memory accumulates continuously
Multi-channel customer or team bot	Hermes	14-platform Gateway with shared memory
One-off tasks, quick answers	Any	Hermes setup overhead not worth it for one-offs
Long-term content series	Hermes + Claude	Hermes for research accumulation, Claude for writing

Cost Comparison

Tool	Monthly Cost	Notes
Claude Pro	$20	Subscription, usage limits apply
Claude Max	$200	High-volume subscription
Hermes (light use)	~$5–10	VPS + API costs; scales with actual usage

Skills Are Portable Across Tools

Hermes follows the **agentskills.io** standard, also used by Claude Code, Cursor, Copilot, and others. Skills you build in one tool work in another. Your accumulated workflows are not locked to any platform.

07 The Honest Caveats

Self-improvement sounds like magic. It's not. Here's what to watch.

Self-Improvement Requires Your Feedback

The learning loop only works as well as the signal it receives. If you use Hermes and never give it feedback — never correct it, never say "this was wrong" or "do it this way instead" — the Skills don't improve meaningfully. The first few weeks are an investment. The payoff is that after several weeks of explicit feedback, it requires much less guidance.

Memory Grows Without Pruning

Hermes doesn't automatically expire old memories. Early mistakes or misunderstandings can persist and subtly affect later behavior if you don't periodically audit what it remembers.

PRACTICAL FIX — Every month or two, spend 20 minutes reviewing the Skills directory and persistent memory. Delete anything wrong or outdated. This is a known area for improvement in future versions.

Autonomous Doesn't Mean Set-and-Forget

Hermes is excellent at optimizing *how* it does things. It's poor at knowing when the *direction* has changed. If your strategy shifts, your product changes, or your standards evolve — Hermes won't pick that up from context alone. You need to tell it, update the relevant Skills, and clear outdated memories.

Database Access Is Real Access

If you connect Hermes to a live database or production system, treat it like any other credential: minimum permissions, read-only where possible. Default MCP configurations often grant more access than necessary. Review and restrict before deploying on anything critical.

Open Source Means You Own the Risk

Hermes is MIT licensed. That means no vendor lock-in, full transparency, and no one can change the terms on you. It also means Nous Research has no commercial obligation to fix something that breaks for you. If you deploy this in a business-critical context, have someone technical who can troubleshoot, or treat it as a workflow enhancer rather than a mission-critical dependency.

Getting Started: The Minimum Viable Path

If you want to try Hermes without over-investing, here's the shortest path to a working setup that delivers real value.

Day 1	• Get a $5 VPS (Hetzner or DigitalOcean, Ubuntu 22.04) • Run the one-line install script • Add an API key (OpenRouter is the simplest starting point) • Create a Telegram bot via @BotFather, add the token to config • Launch Hermes, send it a message from your phone
Week 1	• Use it daily for tasks you'd normally do in a chat window • Give explicit feedback when output isn't right ("do it this way instead") • Watch the Skills directory start to populate automatically
Week 3	• Set up one scheduled task (daily briefing or weekly report) • Connect one MCP integration you actually use (GitHub is usually first)
Month 2+	• The learning loop is running and compounding • Hermes is meaningfully more useful than week one — without additional setup